



# Hands-on with MPI Programming and Spartan

Farzad Khodadadi  
PhD Candidate @ Melbourne eResearch Group (MeG)  
The University of Melbourne  
[farzad.khodadadi@unimelb.edu.au](mailto:farzad.khodadadi@unimelb.edu.au)

# Outline

- MPI programming basics and common guidelines
- MPI4Py demo
- Accessing Spartan, submitting jobs, and monitoring the results

# MPI Programming Basics

Many parallel programs can be written using just these six functions:

```
MPI_INIT  
MPI_FINALIZE  
MPI_COMM_SIZE  
MPI_COMM_RANK  
MPI_SEND  
MPI_RECV
```

- MPI\_SEND and MPI\_RECV functions can be substituted with collective operations such as MPI\_BCAST and MPI\_REDUCE

# Collective Operations in MPI

- ★ **MPI\_BCAST** distributes data from one process (the root) to all others in a communicator.
- ★ **MPI\_REDUCE** combines data from all processes in communicator and returns it to one process.
- ★ In many numerical algorithms, **SEND/RECEIVE** can be replaced by **BCAST/REDUCE**, improving both simplicity and efficiency.

# MPI4Py Sample Programs

```
from mpi4py import MPI  
import sys
```

```
size = MPI.COMM_WORLD.Get_size()  
rank = MPI.COMM_WORLD.Get_rank()  
print("Helloworld! I am process %d of %d.\n" % (rank,  
size))
```

```

import numpy as np
from mpi4py import MPI

from parutils import pprint

comm = MPI.COMM_WORLD

pprint("-"*78)
pprint(" Running on %d cores" % comm.size)
pprint("-"*78)

comm.Barrier()

# Prepare a vector of N=5 elements to be broadcasted...
N = 5
if comm.rank == 0:
    A = np.arange(N, dtype=np.float64)    # rank 0 has proper data
else:
    A = np.empty(N, dtype=np.float64)    # all other just an empty array

# Broadcast A from rank 0 to everybody
comm.Bcast( [A, MPI.DOUBLE] )

# Everybody should now have the same...
print "[%02d] %s" % (comm.rank, A)

```

```

import numpy as np
from mpi4py import MPI

from parutils import pprint

comm = MPI.COMM_WORLD

pprint("-"*78)
pprint(" Running on %d cores" % comm.size)
pprint("-"*78)

my_N = 4
N = my_N * comm.size

if comm.rank == 0:
    A = np.arange(N, dtype=np.float64)
else:
    A = np.empty(N, dtype=np.float64)

my_A = np.empty(my_N, dtype=np.float64)

# Scatter data into my_A arrays
comm.Scatter( [A, MPI.DOUBLE], [my_A, MPI.DOUBLE] )

pprint("After Scatter:")
for r in xrange(comm.size):
    if comm.rank == r:
        print "[%d] %s" % (comm.rank, my_A)
    comm.Barrier()

# Everybody is multiplying by 2
my_A *= 2

# Allgather data into A again
comm.Allgather( [my_A, MPI.DOUBLE], [A, MPI.DOUBLE] )

pprint("After Allgather:")
for r in xrange(comm.size):
    if comm.rank == r:
        print "[%d] %s" % (comm.rank, A)
    comm.Barrier()

```

# MPI4Py Demo





# Parallel Programming using Spartan

- Login to Spartan
  - `yourusername@spartan.hpc.unimelb.edu.au`
- Upload your data (for your assignment, the data has been already uploaded and you only need to create a symbolic link to it in your home directory)
- Write a script to automate execution of your tasks
- Use SLURM's commands to submit your script, monitor your job's execution, cancel it, and much more.

# Spartan Demo



# More on SLURM and it's commands?

- [https://dashboard.hpc.unimelb.edu.au/getting\\_started/](https://dashboard.hpc.unimelb.edu.au/getting_started/)
- <https://rc.fas.harvard.edu/resources/documentation/convenient-slurm-commands/>

[illegible]